

Pseudonymized but Linkable: Cross-Document Re-identification Risks in LLM-Ready Sensitive Text Corpora

Anonymous authors

Paper under double-blind review

Abstract

1 Sensitive text corpora are commonly transformed one document at a time
2 before retrieval, evaluation, sharing, or model adaptation. This work-
3 flow removes obvious direct identifiers, but it can leave a corpus-level
4 privacy failure mode: records with no names, emails, or account numbers
5 may remain linkable through repeated quasi-identifiers such as role, re-
6 gion, schedule, institution type, family context, health context, financial
7 context, and rare events. We introduce CrossDoc-PrivacyBench, a fully
8 synthetic benchmark with 120 controlled personas, 480 documents, risk
9 tiers, utility labels, and decoy auxiliary profiles. We evaluate document-
10 local transformations under linkage, auxiliary-profile matching, RAG-style
11 exposure, quasi-identifier recovery, and utility metrics. Direct redaction
12 leaves auxiliary-profile matching unchanged at 0.71 top-1, while a Presidio-
13 style PII baseline remains highly matchable at 0.59. We propose Link-
14 Guard, a simple corpus-aware generalization heuristic that estimates cross-
15 document uniqueness and generalizes high-risk quasi-identifiers until a
16 target anonymity set is reached. At target $k = 5$, LinkGuard reduces
17 auxiliary top-1 to 0.04 while preserving issue classification and retrieval
18 Recall@5 at 1.00; aggressive redaction reaches similar privacy but drops
19 issue accuracy to 0.35 and retrieval Recall@5 to 0.31. GPT-5.5 auxiliary and
20 RAG-generation stress audits preserve this ordering, reinforcing that LLM-
21 ready data transformations should be evaluated for corpus-level linkage
22 resistance, not only document-local PII removal.

23 1 Introduction

24 Sensitive data is often most useful when it is not isolated. A hospital may have clinical
25 messages, billing notes, scheduling records, and accommodation requests about the same
26 patient. A legal-aid organization may have intake notes, benefits appeals, landlord dis-
27 putes, and follow-up emails about the same client. A company may have support tickets,
28 HR notes, finance requests, and internal chat snippets about the same employee or cus-
29 tomer. Foundation-model workflows, retrieval-augmented generation (RAG), evaluation-set
30 construction, model adaptation, analytics, and data sharing, therefore often operate over
31 *multi-document* corpora rather than one-off documents.

32 The dominant operational response is document-local de-identification. A pipeline removes
33 names, emails, phone numbers, mailing addresses, account numbers, and sometimes other
34 spans that look individually identifying. This is necessary, but it is not sufficient. A
35 rare combination such as a region, occupation, institution type, care schedule, family
36 constraint, legal context, and local event may be harmless in one document yet become
37 a stable signature when repeated across documents. Pseudonymization can make this
38 worse: a stable pseudonym is useful for continuity, but it is also a linkage handle. Per-
39 document pseudonyms remove the explicit handle, but they do not remove the repeated
40 quasi-identifier structure.

41 This paper studies a focused question: *after document-local anonymization or pseudonymization,*
42 *can an attacker still link multiple transformed documents to the same person and match them to aux-*

43 *iliary background profiles?* We intentionally avoid real-person data and real re-identification
 44 attempts. Instead, we build a controlled synthetic benchmark in which every persona,
 45 document, decoy, and utility label is generated synthetically to give us safe ground truth for
 46 controlled privacy auditing.

47 Our contributions are:

- 48 • **Benchmark:** CrossDoc-PrivacyBench, a synthetic cross-document privacy bench-
 49 mark with 120 personas, 480 documents, risk tiers, utility labels, templates, and
 50 decoy auxiliary profiles.
- 51 • **Audit protocol:** pairwise linkage, fixed- K clustering, auxiliary-profile matching,
 52 quasi-identifier recovery, RAG-style exposure, and utility retention.
- 53 • **Finding:** direct redaction, Presidio-style PII removal, document-local anonymiza-
 54 tion, and stable pseudonyms all leave strong residual linkage risk.
- 55 • **Defense:** LinkGuard, a simple corpus-aware generalization heuristic that improves
 56 the synthetic privacy–utility tradeoff over blanket aggressive redaction.

57 2 Problem Setup and Threat Model

58 **Setting.** A data holder has a sensitive text corpus with multiple records per person,
 59 household, customer, patient, client, or employee. The corpus is transformed before it is
 60 indexed for RAG, shared internally, released for evaluation, or used in model-adaptation
 61 workflows. We consider transformations that operate on text, including direct redaction,
 62 pseudonymization, document-local anonymization, and corpus-aware generalization.

63 **Attacker knowledge.** The attacker sees the transformed corpus. For the auxiliary-profile
 64 task, the attacker also receives a small candidate set containing one true synthetic auxiliary
 65 profile and decoys. This simulates background knowledge such as an external directory,
 66 leaked structured metadata, public biographical context, or another internal table. The at-
 67 tacker does not receive the original direct identifiers, and the candidate labels are neutralized
 68 to avoid leaking synthetic persona IDs.

69 **Attacker goals.** We evaluate three defensive questions. First, can an attacker tell which
 70 transformed documents belong to the same synthetic persona? Second, can the attacker
 71 match a cluster of transformed documents to the correct auxiliary profile among decoys?
 72 Third, can profile-like queries retrieve target documents from a transformed corpus, creating
 73 a RAG exposure channel?

74 **Non-goals and ethics.** We do not collect, process, or attempt to re-identify real peo-
 75 ple. All personas, documents, auxiliary profiles, and decoys are synthetic. The bench-
 76 mark is intended for defensive auditing of transformation methods, not for operational
 77 re-identification. Any extension to real data would require authorization, data minimization,
 78 security review, and domain-specific ethical oversight.

79 3 CrossDoc-PrivacyBench: A Synthetic Cross-Document Benchmark

80 **Personas.** CrossDoc-PrivacyBench contains 120 synthetic personas. Each persona has
 81 direct identifiers: name, email, phone number, address, and account ID, and a structured
 82 set of quasi-identifiers: region, city, occupation, employer type, education, family structure,
 83 medical context, financial context, legal context, schedule pattern, community affiliation,
 84 and rare event. Each persona also has utility labels for four domains. We use deterministic
 85 generation so that experiments are reproducible and no real individual is represented.

86 **Risk tiers.** Personas are assigned to three tiers. T1 documents contain mostly generic
 87 quasi-identifiers, so direct PII is the main privacy risk. T2 documents contain moderately
 88 specific but partly generalized context, such as a city within a region, employer category,
 89 broad health or legal context, and schedule constraints. T3 documents contain high-linkage

90 combinations, such as exact city, occupation, employer type, education, family structure,
 91 treatment or schedule context, affiliation, and rare event. This tiering lets us check whether
 92 an attack and defense are sensitive to increasing quasi-identifier specificity.

93 **Documents.** Each persona has four documents spanning healthcare, legal, financial sup-
 94 port, and HR. The documents are short service or support records with a direct-contact
 95 line and a task-oriented body. Each document includes an issue label, such as medication
 96 refill, benefits appeal, payment deferral, housing issue, accommodation request, or sched-
 97 ule change. The benchmark therefore contains 480 documents total and 384 held-out test
 98 documents from 96 held-out personas.

99 **Auxiliary profiles and decoys.** For each persona, we construct a ten-candidate auxiliary
 100 set. One candidate is the true profile; nine are decoys selected to share some attributes
 101 with the target, such as region, occupation family, or legal context. This makes matching
 102 harder than a trivial exact lookup and lets us report top-1, top-3, and mean reciprocal rank.
 103 The auxiliary-profile task is the central privacy test because it captures both linkage and
 104 background-knowledge use.

105 4 Transformations

106 We compare eight conditions. C0 is the original synthetic corpus. C1 performs direct
 107 redaction of names, first and last names, email addresses, phone numbers, mailing addresses,
 108 and account IDs. C1b applies a Presidio-based PII analyzer plus local regex cleanup for
 109 the synthetic account and address formats. C2 uses stable pseudonyms across the corpus,
 110 replacing direct identifiers and some locations/institutions with consistent tokens. C3 resets
 111 pseudonym mappings per document. C4 is a document-local anonymization proxy that
 112 removes direct identifiers and generalizes selected local spans, such as city and education,
 113 without seeing neighboring documents. C5 is LinkGuard. C6 is aggressive redaction, which
 114 suppresses most potentially identifying context, including quasi-identifiers and task issue
 115 phrases.

116 4.1 LinkGuard: Corpus-Aware Generalization

117 Document-local anonymizers ask whether a span is identifying inside one document. Link-
 118 Guard asks whether repeated details make a persona distinguishable in the corpus. It
 119 operates over persona groups in the synthetic benchmark, but the same logic could be
 120 applied to known customer, patient, case, or account groupings in a controlled deployment.

For each persona and quasi-identifier field, LinkGuard defines a three-level transformation ladder:

exact value $\rightarrow$ semantic category $\rightarrow$ generic suppression.

For example, a city can become a region and then “a broad location”; an occupation can become a job family and then “general worker”; a rare event can become a local-event category and then “a local event.” The method starts from direct redaction and estimates how many personas share the current generalized representation. If the estimated anonymity set is smaller than the target k , LinkGuard selects another field to generalize. The selection score balances rarity, a specificity prior, and an approximate utility cost:

$$s(f) = \frac{\text{rarity}(f) + \text{specificity}(f) + 1}{\max(\text{utility_cost}(f), 0.1)}.$$

121 Fields with rarer values, higher specificity, and lower expected utility cost are generalized
 122 first. The process stops when the estimated anonymity set reaches target k or all fields
 123 have been generalized or suppressed. This is a heuristic risk-budgeting procedure, the
 124 estimate depends on the modeled population, the field schema, and the attacker’s auxiliary
 125 knowledge.

5 Attacks and Metrics

Pairwise linkage. We compute signature features over transformed documents, including explicit pseudonym handles and quasi-identifier phrases. Pairwise linkage is evaluated as same-person binary classification over held-out document pairs. We select a threshold on validation personas and report test F1. This metric is especially useful for showing that stable pseudonyms can become linkage handles.

Fixed- K clustering. We also report fixed- K clustering adjusted Rand index, where K is the true number of held-out personas. Threshold graph clustering is brittle outside stable-handle settings, so we treat pairwise linkage and auxiliary matching as the main privacy evidence.

Auxiliary-profile matching. For each held-out persona, the attacker ranks the ten auxiliary candidates using the transformed documents. The main local attacker uses word TF-IDF similarity; robustness checks use character TF-IDF, hybrid TF-IDF, and a field-weighted matcher that explicitly scans exact and generalized schema fields. We report top-1, top-3, and mean reciprocal rank. Lower values indicate better privacy.

Exact quasi-identifier recovery. Auxiliary matching gives an end-to-end privacy signal, but it can hide which information leaked. We therefore scan transformed documents for exact recovery of quasi-identifier fields such as city, occupation, employer type, education, family context, medical context, financial context, legal context, schedule pattern, affiliation, and rare event. This is a structured profile-reconstruction diagnostic.

RAG exposure. To model retrieval exposure, we issue synthetic profile queries against the transformed corpus and measure whether target-persona documents appear in the top retrieved results. We report Hit@5, whether multiple target documents appear in the top ten, document recall at ten, and reciprocal rank. This diagnostic asks whether transformed records remain reachable from profile-like background knowledge even when they contain no direct identifiers.

Utility. We evaluate lightweight task utility with issue classification, task-oriented retrieval Recall@5, and fact preservation. The classifiers and retrievers are local TF-IDF models trained on original synthetic documents. These metrics test whether service-triage information survives transformation; they do not claim to cover all downstream uses.

6 Results

6.1 Main Privacy–Utility Sweep

Table 1 summarizes the main local sweep over 96 held-out personas. Direct redaction leaves auxiliary matching unchanged at 0.7 top-1, confirming that direct PII removal alone does not remove repeated quasi-identifier risk. The Presidio-style baseline improves over direct redaction but remains highly matchable at 0.6. The document-local proxy still leaves auxiliary top-1 at 0.6 and exact quasi-identifier recovery at 0.3. Consistent pseudonymization makes linkage nearly trivial, with pairwise F1 of 0.98, because the stable token itself becomes a cross-document handle.

LinkGuard reduces auxiliary top-1 to 0.04 and exact recovery to 0.0 while preserving issue accuracy and retrieval Recall@5 at 1.0. Aggressive redaction reaches a similar auxiliary top-1 of 0.05, but issue accuracy drops to 0.35 and retrieval Recall@5 to 0.31. Thus, in this synthetic sprint, corpus-aware generalization occupies a better privacy–utility point than blanket suppression. A stricter body-only utility stress test tells the same story: LinkGuard scores 0.88, direct redaction 0.87, and aggressive redaction 0.44.

Cond.	Pair	Aux@1	Exact	Issue	Ret@5	Facts	Edit
C0 Orig	0.96	0.71	0.39	1.00	1.00	0.68	0.00
C1 Redact	0.21	0.71	0.39	1.00	1.00	0.68	0.25
C1b Presidio	0.14	0.59	0.33	1.00	1.00	0.68	0.35
C2 Stable	0.98	0.43	0.27	1.00	1.00	0.68	0.30
C3 PerDoc	0.15	0.43	0.27	1.00	1.00	0.68	0.31
C4 Local	0.16	0.60	0.30	1.00	1.00	0.68	0.29
C5 LinkGuard	0.02	0.04	0.00	1.00	1.00	0.70	0.54
C6 Agg	0.02	0.05	0.00	0.35	0.31	0.39	0.68

Table 1: Main held-out results. Pair is pairwise linkage F1; Aux@1 is auxiliary-profile top-1 accuracy; Exact is exact quasi-identifier recovery; Issue and Ret@5 are utility metrics; Facts is lightweight fact preservation; Edit is character-level edit ratio. Lower Pair, Aux@1, Exact, and Edit are better for privacy or fidelity; higher Issue, Ret@5, and Facts are better for utility.

Cond.	T1 Aux@1	T2 Aux@1	T3 Aux@1	T3 Exact
C1 Redact	0.16	0.97	1.00	1.00
C1b Presidio	0.09	0.69	1.00	0.88
C4 Local	0.16	0.66	1.00	0.82
C5 LinkGuard	0.03	0.00	0.09	0.00
C6 Agg	0.06	0.00	0.09	0.00

Table 2: Risk-tier stratification. Direct and document-local baselines fail most strongly on high-linkage T3 personas; LinkGuard reduces both auxiliary matching and exact quasi-identifier recovery.

6.2 Risk Tiers and Signal Ablations

The benchmark tiers behave as intended. Under direct redaction, Aux@1 rises from 0.16 on T1 to 0.97 on T2 and 1.00 on T3. LinkGuard lowers the corresponding rates to 0.03, 0.00, and 0.09, and reduces T3 exact quasi-identifier recovery to 0.00. The result indicates that direct redaction is adequate only in the tier where quasi-identifiers were intentionally generic; it fails as soon as the corpus contains repeated, moderately specific context.

Single-signal ablations of direct-redacted documents indicate that role and institution are the largest local drivers. Removing role and organization drops Aux@1 from 0.71 to 0.50, while removing location drops it to 0.52. Family, health, rare event, and schedule ablations have little effect individually in this synthetic corpus, but they still matter in combination and under the field-weighted stress attacker. This is consistent with the central claim: linkage risk is often combinatorial, not span-local.

6.3 Sensitivity to the Linkage Risk Budget

Table 3 varies the target anonymity set size used by LinkGuard. Moving from direct redaction ($k = 1$) to $k = 2$ drops Aux@1 from 0.71 to 0.06 and exact recovery from 0.39 to 0.01; $k = 5$ removes exact recovery and lowers the field-weighted attacker to 0.24. Larger k mainly helps the field-weighted stress attacker, which falls to 0.10 at $k = 20$, so corpus-aware generalization acts as a tunable risk budget. A corpus-awareness ablation confirms that true corpus statistics matter: target- k coverage is 1.00 with true planning but 0.92 under shuffled quasi-identifier co-occurrences, while a global level-1 rule leaves field-aware top-1 at 0.98.

6.4 Attack Robustness, RAG Exposure, and API Audit

Across word, character, hybrid TF-IDF, and field-weighted auxiliary matchers, direct redaction ranges from 0.51 to 0.98 top-1 and the document-local proxy ranges from 0.48 to 0.98. LinkGuard remains lower, from 0.03 to 0.24, but the field-weighted result exposes residual structured-context risk. The residual analysis finds four LinkGuard top-1 residual matches under the main word attacker; none retain exact quasi-identifier fields, and the median

k	Min k	Edit	Pair	Aux@1	Field@1	Exact	Issue	Ret@5
1	1	0.25	0.21	0.71	0.98	0.39	1.00	1.00
2	2	0.52	0.04	0.06	0.35	0.01	1.00	1.00
3	3	0.53	0.03	0.05	0.33	0.00	1.00	1.00
5	5	0.54	0.02	0.04	0.24	0.00	1.00	1.00
8	8	0.54	0.02	0.04	0.17	0.00	1.00	1.00
12	12	0.54	0.02	0.04	0.12	0.00	1.00	1.00
20	20	0.54	0.02	0.04	0.10	0.00	1.00	1.00

Table 3: LinkGuard sensitivity to the target anonymity set size. The main auxiliary matcher drops sharply by $k = 2$, while the field-weighted stress attacker continues improving at larger k .

Table 4: GPT-5.5 qualitative evidence extraction over 24 selected synthetic cases. Entries are case-level rates for signal families in the extracted explanations.

Case	n	Role	Loc.	Inst.	High-spec.	Unc.
Direct success	8	0.75	1.00	0.50	0.47	0.00
LG residual	8	0.00	0.00	0.00	0.07	0.88
Agg failure	8	0.00	0.00	0.00	0.18	0.75

score margin is only 0.00. This means that LinkGuard does not eliminate all matching signal, but the remaining matches are weak and mostly coarse.

The RAG-style exposure diagnostic shows why linkage evaluation matters for LLM-ready corpora. For high-linkage T3 personas, exact profile queries retrieve a target document for every direct-redacted, Presidio-redacted, and document-local corpus; LinkGuard reduces T3 Hit@5 to 0.31 and Multi@10 to 0.00, while aggressive redaction reduces Hit@5 to 0.03 at much higher utility cost. Deterministic generated profile-like queries preserve the ordering: verbose-query Hit@5 is 0.38 for direct redaction and the document-local proxy, 0.26 for LinkGuard, and 0.03 for aggressive redaction; short role-region queries separate direct redaction from LinkGuard at 0.35 versus 0.02. A deterministic top-5 RAG context-recovery scan on T3 profiles shows that direct redaction exposes 9.78 exact quasi-identifier fields on average and the document-local proxy exposes 7.84 exact / 9.47 coarse fields, while LinkGuard exposes 0.00 exact / 0.44 coarse fields. This diagnostic is not a full RAG attack, but it shows that transformed corpora can remain searchable by auxiliary profiles.

Finally, time-stamped GPT-5.5 stress audits use synthetic data only, capped calls, cached responses, and store=false. Auxiliary matching preserves the ordering: direct, Presidio, and document-local transformations score 1.00, 0.98, and 1.00 top-1, while LinkGuard drops to 0.44 with uncertainty 0.92 and aggressive redaction scores 0.33. A 24-person GPT-5.5 document-local anonymization baseline remains fully matchable at 1.00 top-1. An additional 24-case GPT-5.5 evidence-extraction audit finds that direct-redaction successful matches cite location in 1.00 and role in 0.75 of cases, while LinkGuard residual matches cite role, location, and institution at 0.00 and are uncertain in 0.88. A completed 12-person T3 RAG-generation audit labels direct/Presidio/document-local contexts as same-person at 1.00, versus 0.42 uncertain for LinkGuard and 0.00 for aggressive redaction. Because API model behavior can change, these are stress audits rather than the main statistical evidence.

7 Related Work

Classic work on k -anonymity formalized how released records can remain re-identifiable through combinations of quasi-identifiers and motivated generalization and suppression (Sweeney, 2002). Subsequent privacy models such as ℓ -diversity and t -closeness highlighted that anonymity-set size alone is insufficient when sensitive attributes remain homogeneous or distributionally revealing (Machanavajjhala et al., 2007; Li et al., 2007). Sparse-data attacks showed that high-dimensional records can be de-anonymized with limited auxiliary knowledge (Narayanan and Shmatikov, 2008), and later work examined how practical

de-identification defenses can fail under background knowledge and downcoding attacks (Cohen, 2022).

Text anonymization benchmarks such as TAB evaluate anonymization decisions and utility in legal text (Pilán et al., 2022). RAT-Bench recently reframes text anonymization around re-identification risk from both direct and indirect identifiers (Krčo et al., 2026). Surveys of text anonymization and anonymization metrics emphasize that privacy, utility, legal requirements, and user-centered expectations are often measured inconsistently (Deußer et al., 2025; Ren et al., 2025). Presidio and OpenAI Privacy Filter illustrate the practical importance of PII detection tools, but such tools are not by themselves complete anonymization guarantees (Microsoft, 2026; OpenAI, 2026).

Recent work on LLM privacy shows that language models can infer personal attributes from text beyond direct memorization (Staab et al., 2024), and that large models can memorize and reveal training data under extraction attacks (Carlini et al., 2021; Nasr et al., 2023). RAG privacy work shows that retrieval databases can leak private records through retrieval and generation channels (Zeng et al., 2024). Autoregressive infilling attacks on de-identified documents further show that background knowledge can help recover masked spans (Charpentier and Lison, 2025). Our contribution is complementary: we focus on cross-document linkage after document-level transformations, where repeated corpus-level quasi-identifiers can survive otherwise reasonable LLM-ready anonymization workflows.

8 Limitations and Ethics

The benchmark is synthetic and template-controlled. This is necessary for safe ground truth and reproducibility, but it limits external validity. Real sensitive corpora may contain richer writing styles, temporal dependencies, institutional terminology, missing fields, contradictions, OCR errors, and adversary knowledge. Our local attacks are lightweight, and the GPT-5.5 audit is a time-stamped API stress test rather than a stable benchmark. LinkGuard is a heuristic generalization method, not a formal privacy guarantee. The field-weighted attacker shows that residual generalized context can still be informative. Utility is measured with issue labels, retrieval, and lightweight fact checks rather than full downstream deployment tasks.

The benchmark is intended for defensive auditing. We do not collect, process, or attempt to re-identify real people. All examples are synthetic and generated from a controlled schema. Any extension to real data would require authorization, data minimization, access controls, logging, and ethical review. Practitioners should not treat successful performance on this benchmark as proof that a real corpus is safe to release or train on.

9 Conclusion

Document-local PII removal is not enough for LLM-ready sensitive corpora. In a multi-document setting, repeated quasi-identifiers can support linkage, auxiliary-profile matching, and profile-query retrieval even after direct identifiers are removed. Our synthetic benchmark suggests that privacy evaluation should include corpus-level linkage resistance, not only span-level PII recall. It also suggests that simple corpus-aware generalization can improve the privacy-utility frontier relative to document-local anonymization and blanket redaction. The broader lesson is that data transformation for foundation models should ask what a capable attacker can still infer across the corpus, not only what identifying spans were removed from each document.

References

Nicholas Carlini, Florian Tramèr, Eric Wallace, Matthew Jagielski, Ariel Herbert-Voss, Katherine Lee, Adam Roberts, Tom Brown, Dawn Song, Úlfar Erlingsson, Alina Oprea, and Colin Raffel. Extracting training data from large language models. In *USENIX Security Symposium*, 2021.

- 279 Lucas Charpentier and Pierre Lison. Re-identification of de-identified documents with
280 autoregressive infilling. *arXiv:2503.03438*, 2025.
- 281 Aloni Cohen. Attacks on deidentification’s defenses. In *USENIX Security Symposium*, 2022.
- 282 Theresa Deußer, Andreas Kirsch, Florian Matthes, and Simon Mayer. SoK: A survey of text
283 anonymization techniques. *arXiv:2508.15433*, 2025.
- 284 Nataša Krčo, Zexi Yao, Matthieu Meeus, and Yves-Alexandre de Montjoye. RAT-Bench: A
285 comprehensive benchmark for text anonymization. *arXiv:2602.12806*, 2026.
- 286 Ninghui Li, Tiancheng Li, and Suresh Venkatasubramanian. t -closeness: Privacy beyond
287 k -anonymity and ℓ -diversity. In *IEEE International Conference on Data Engineering*, 2007.
- 288 Ashwin Machanavajjhala, Daniel Kifer, Johannes Gehrke, and Muthuramakrishnan Venkata-
289 subramaniam. ℓ -diversity: Privacy beyond k -anonymity. *ACM Transactions on Knowledge*
290 *Discovery from Data*, 2007.
- 291 Microsoft. Presidio: Data protection and de-identification SDK. [https://microsoft.github.](https://microsoft.github.io/presidio/)
292 [io/presidio/](https://microsoft.github.io/presidio/), 2026.
- 293 Arvind Narayanan and Vitaly Shmatikov. Robust de-anonymization of large sparse datasets.
294 In *IEEE Symposium on Security and Privacy*, 2008.
- 295 Milad Nasr, Nicholas Carlini, Jonathan Hayase, Matthew Jagielski, A. Feder Cooper, Daphne
296 Ippolito, Christopher A. Choquette-Choo, Eric Wallace, Florian Tramèr, and Katherine Lee.
297 Scalable extraction of training data from production language models. *arXiv:2311.17035*,
298 2023.
- 299 OpenAI. Introducing OpenAI Privacy Filter. [https://openai.com/index/](https://openai.com/index/introducing-openai-privacy-filter/)
300 [introducing-openai-privacy-filter/](https://openai.com/index/introducing-openai-privacy-filter/), 2026.
- 301 Ildikó Pilán, Pierre Lison, Lilja Øvrelid, Anthi Papadopoulou, David Sánchez, and Montser-
302 rat Batet. The Text Anonymization Benchmark (TAB): A dedicated corpus and evaluation
303 framework for text anonymization. *Computational Linguistics*, 2022.
- 304 Hao Ren, Yi Zhang, and Yang Li. A survey of text anonymization metrics. *arXiv:2511.06282*,
305 2025.
- 306 Robin Staab, Mark Vero, Mislav Balunović, and Martin Vechev. Beyond memorization:
307 Violating privacy via inference with large language models. In *International Conference on*
308 *Learning Representations*, 2024.
- 309 Latanya Sweeney. Achieving k -anonymity privacy protection using generalization and
310 suppression. *International Journal of Uncertainty, Fuzziness and Knowledge-Based Systems*,
311 2002.
- 312 Shenglai Zeng, Jiankun Zhang, Pengfei He, Yiding Liu, Yiming Liu, Hanxu Hu, Lianhui
313 Qin, Yue Xing, Mengdi Wang, and Jiliang Tang. The good and the bad: Exploring privacy
314 issues in retrieval-augmented generation. *arXiv:2402.16893*, 2024.